December 11, 2021

The University of Texas at Austin

EVALUATION OF LRU v/s SRRIP CACHE REPLACEMENT POLICY

GRAD LAB PART - 2

Nithyashree Srinivasan (ns32928)
Shreyas Ravishankar (sr48925)
Abhijjith Venkkateshraj (av36677)

Cache Replacement Policies

- Cache Replacement decides which cache line are to be replaced in set associative and fully associative caches.
- Commonly used policies-
  - LRU (Least Recently Used)
  - NRU (Not Recently Used)
  - FIFO
  - Random
  - RRIP

Types of Access patterns and Why LRU doesn’t perform well

- **Recency-friendly access patterns**- Access pattern that repeats N times. LRU works well in this case.
- **Thrashing access patterns**- Cyclic access pattern of length K which repeats N times. K> no. of cache blocks → LRU performs very badly (0 hit rate).
- **Streaming access patterns**- Patterns with infinite re-reference interval. No hits in any replacement policy, so LRU doesn’t matter.
- **Mixed access patterns**- A mix of scan and frequently referenced working set. LRU fails to preserve the working set after a scan.

RRIP

- Based on re-reference interval prediction.
- LRU-MRU block will be re-referenced sooner, and the block at LRU position will be referenced last.
- **Optimal algorithm**- Predicts distance re-reference interval for scan type data and near re-reference interval for working set/frequently used data.
- RRIP prevents scan blocks from replacing the working set
- SRRIP - scan resistant
- DRRIP- scan resistant and thrash resistant

SRRIP (Static RRIP)

- Overhead- M bit register holding RRPV bits (M bit saturating counter) for each cache line.
- Two types
  - HP- Change RRPV = 0 on a hit
  - FP- Change RRPV = RRPV - 1 on a hit
- The figure on the right shows RRIP-HP with a 2 bit saturating counter (M=2)

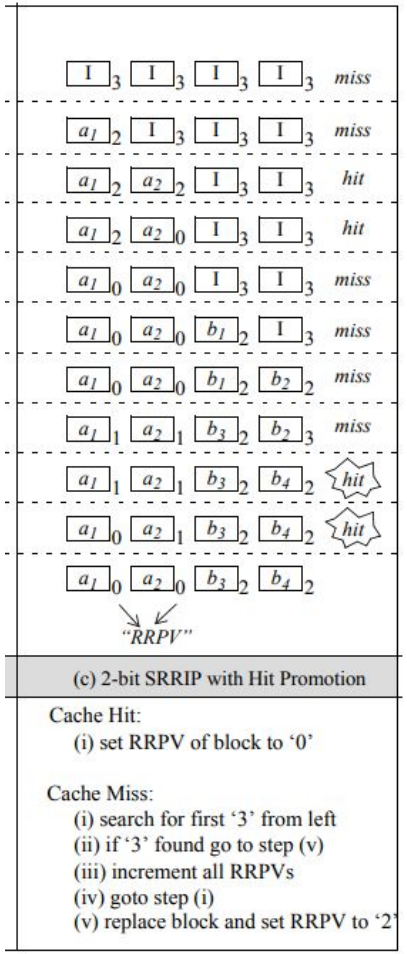

Methodology and Configuration

- Base configuration
  - **1st level**
    - I-Cache 32KB, Associativity 4
    - D-Cache 32KB, Associativity 4
  - **2nd level**
    - LLC 1MB, Associativity 8
  - Prefetcher on
- Evaluation- SRRIP and LRU are compared across various LLC configurations-
  - Fixed Associativity (8) with varying cache sizes 256KB, 512KB, 1MB, 2MB, 4MB, 8MB.
  - Fixed Cache Size (1M) with varying associativity 2,4,8,16.
  - Prefetch vs No-Prefetch with Associativity (16) and Cache Size (1M).

Code Changes

Cache_lib.h

- Cache_Entry_struct- Added rrpv bits
- Added SRRIP replacement to
  - Repl_Policy_enum
  - Cache_Insert_Repl_enum

Cache_lib.c

- Added hit logic for SRRIP to cache_access function.
- Added miss logic for SRRIP find_repl_entry function to find the entry to be replaced.
- Insertion logic is placed in cache_insert_replpos function.

RESULTS

BZIP2- Miss rate variation across associativity

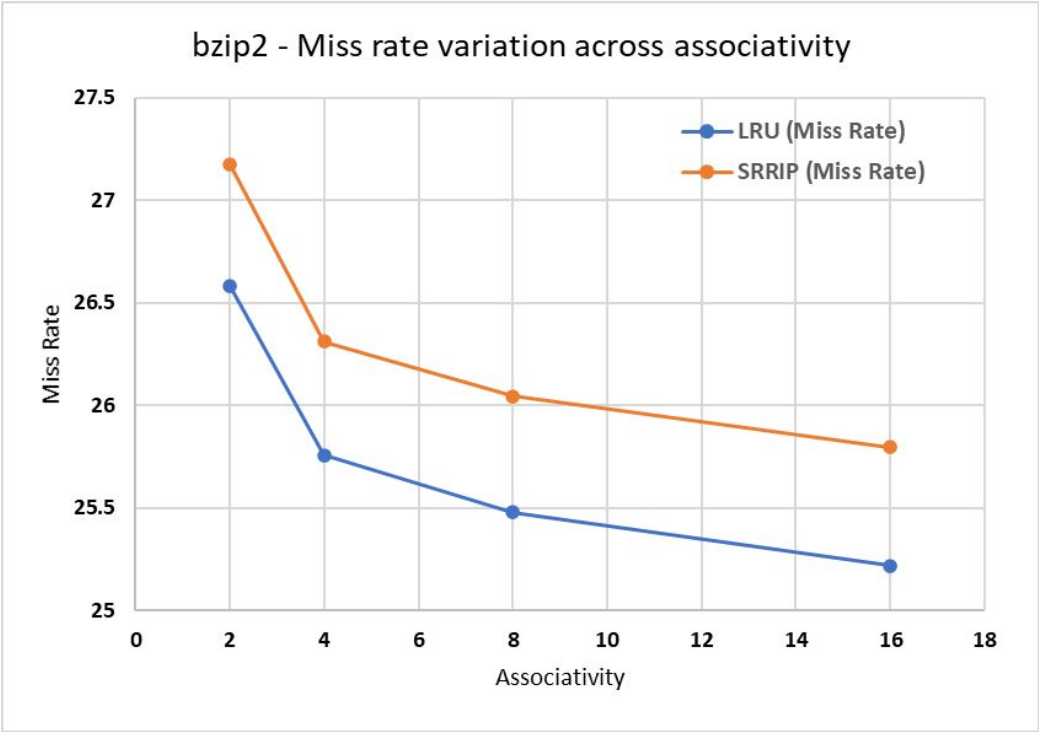

- Miss rate increased from LRU to SRRIP.

HMMR- Miss rate variation across associativity

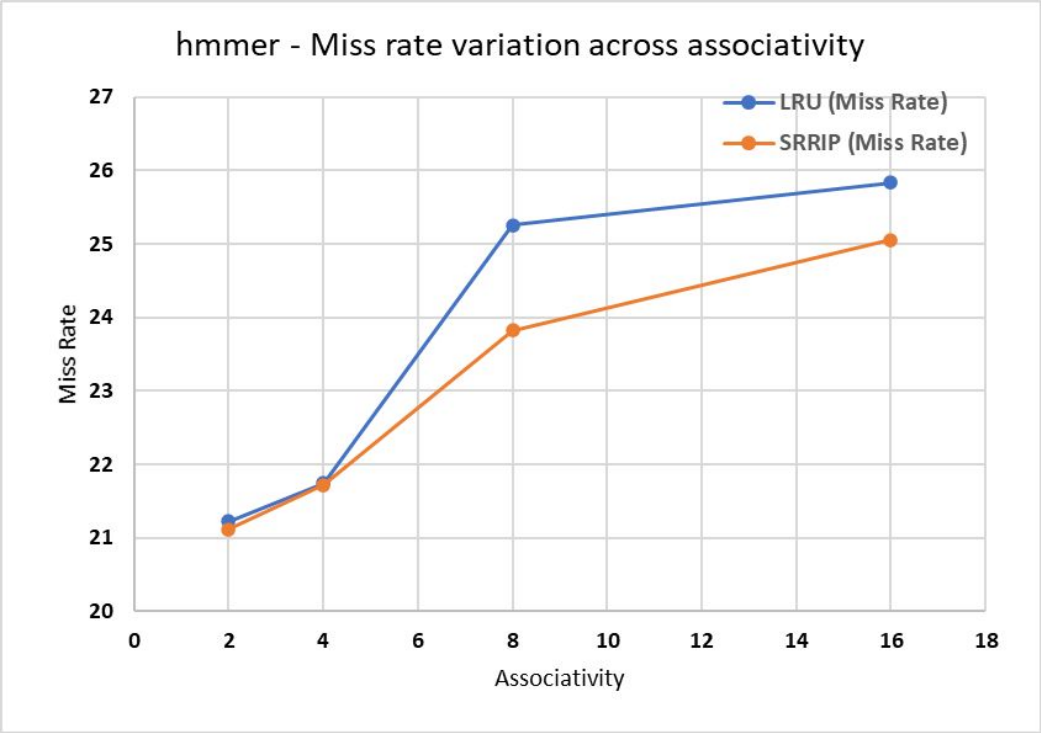

- Miss rate decreased from LRU to SRRIP.

MCF- Miss rate variation across associativity

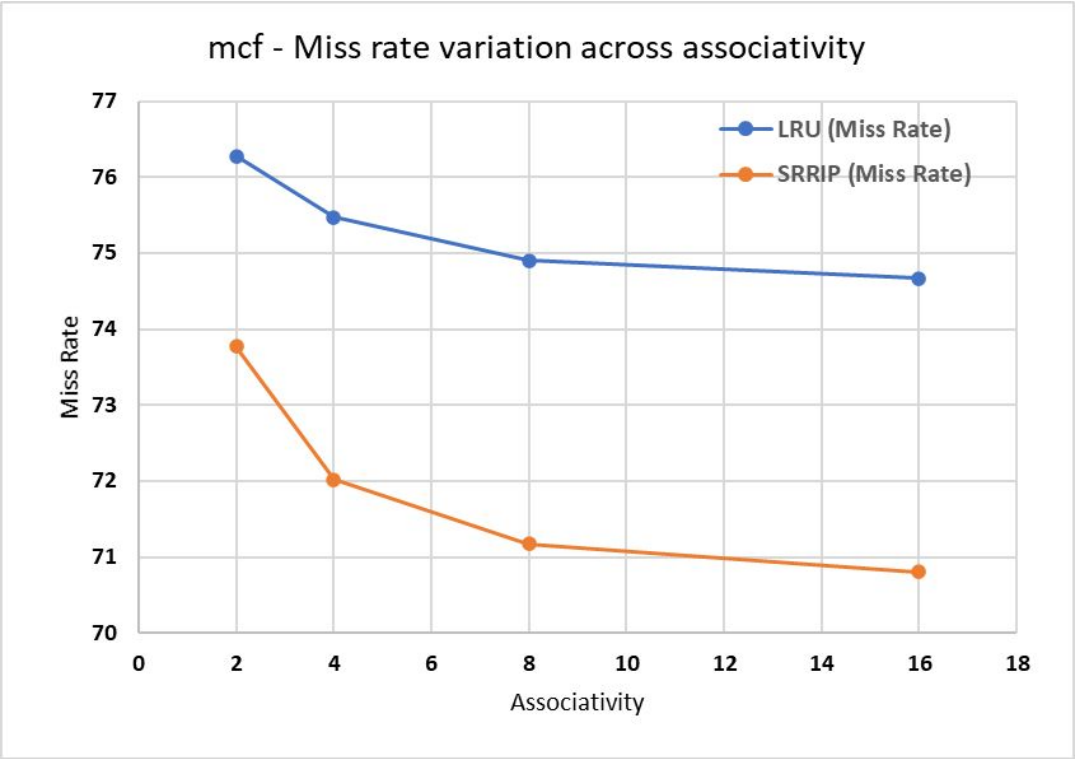

- Significant reduction in the miss rate across all associativities.

Variation of % change in miss rate with associativity

- The graph shows that RRIP doesn't show any benefits for BZIP benchmark.
- The HMMR benchmark shows some benefit in miss rate, which increases with associativity.
- The MCF workload shows the maximum benefits for SRRIP

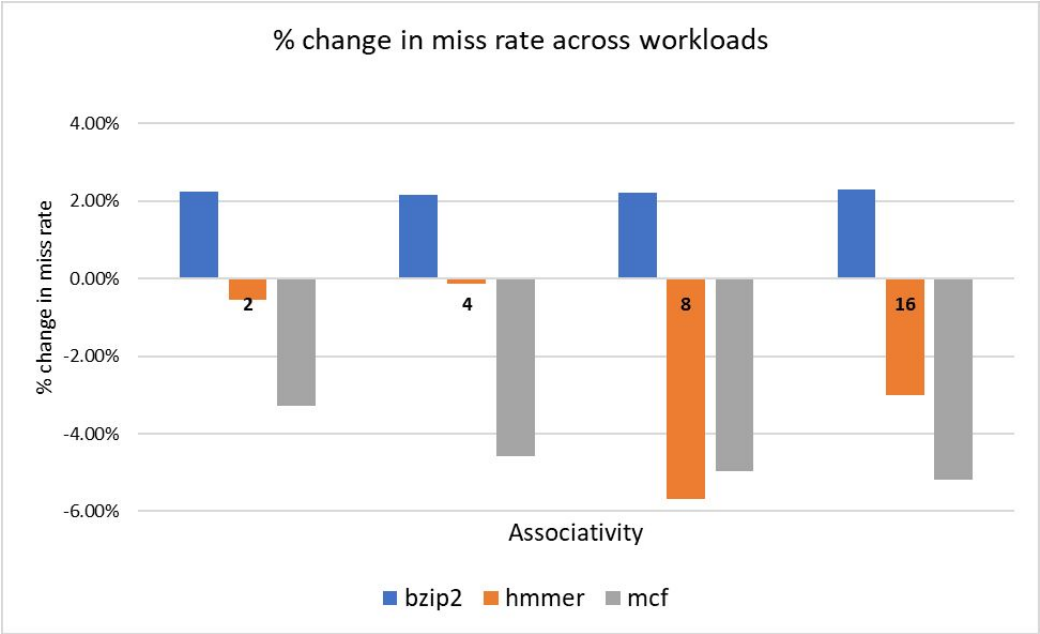

BZIP2- Miss rate variation across memory size

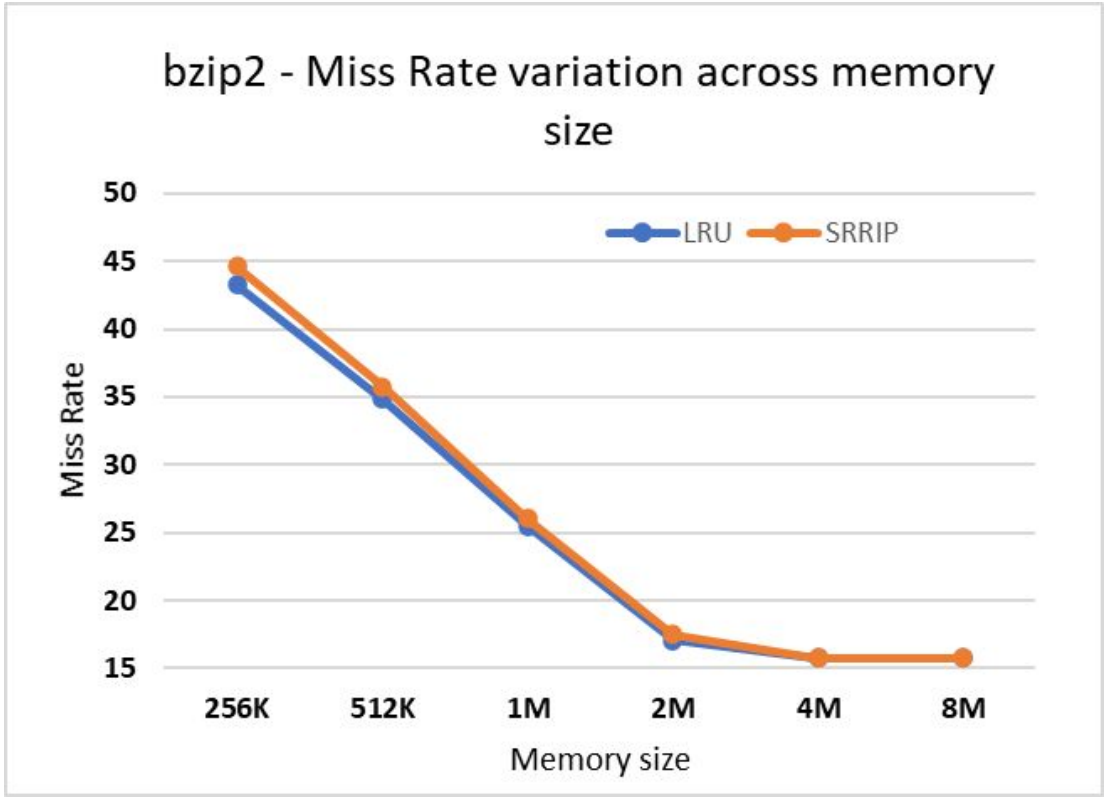

- SRRIP doesn’t show much improvement for this workload.

HMMR- Miss rate variation across memory size

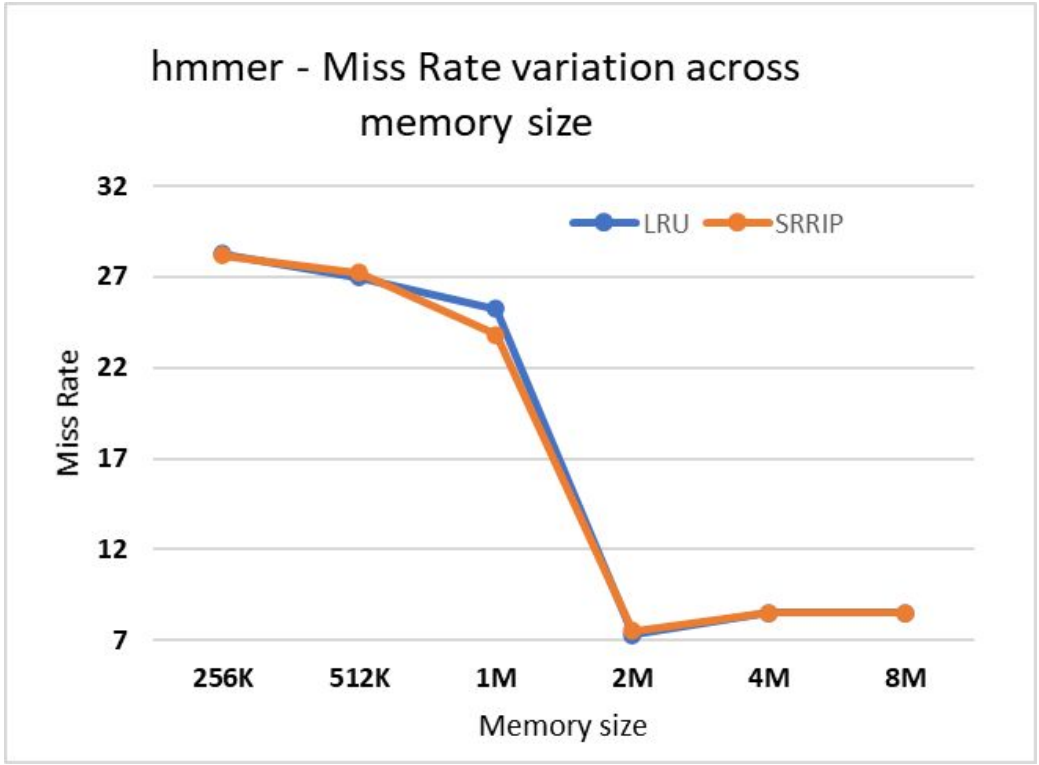

- We don’t see a significant difference in the miss rate.

MCF- Miss rate variation across memory size

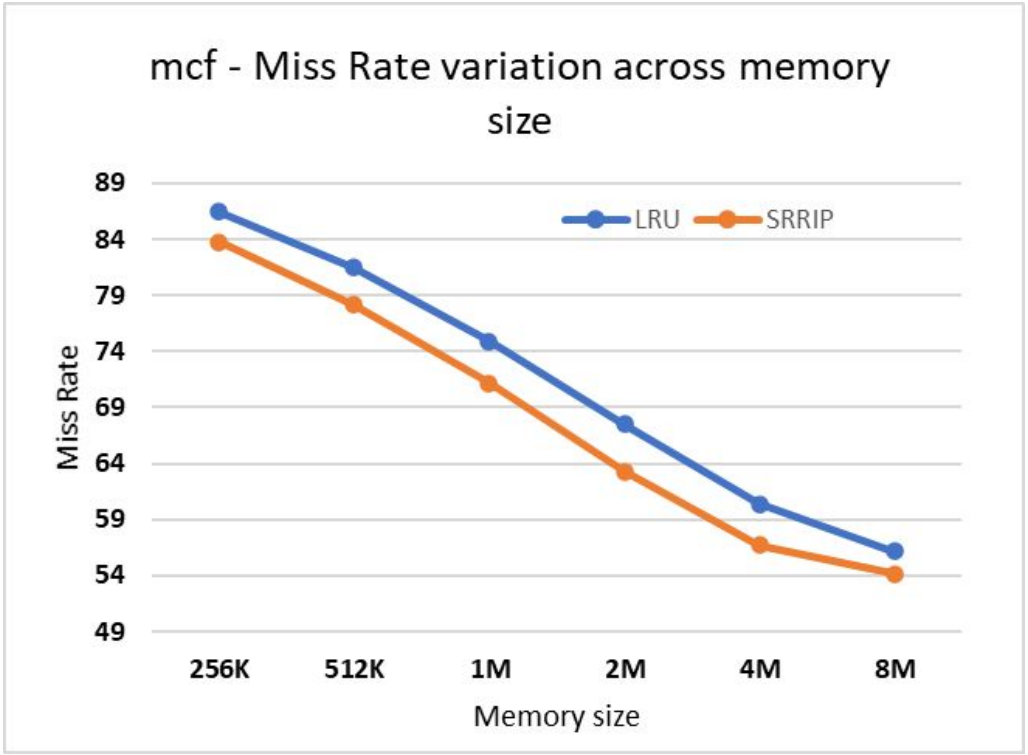

- There is a steady decrease in miss rate for mcf when we vary the memory size.

Variation of % change in miss rate with cache size

- BZIP doesn’t show any improvement across sizes.
- HMMR shows fluctuation with cache size. So the behaviour is erratic.
- MCF shows a reduction in miss rate till 4MB and then the benefits begin to reduce for greater cache sizes.

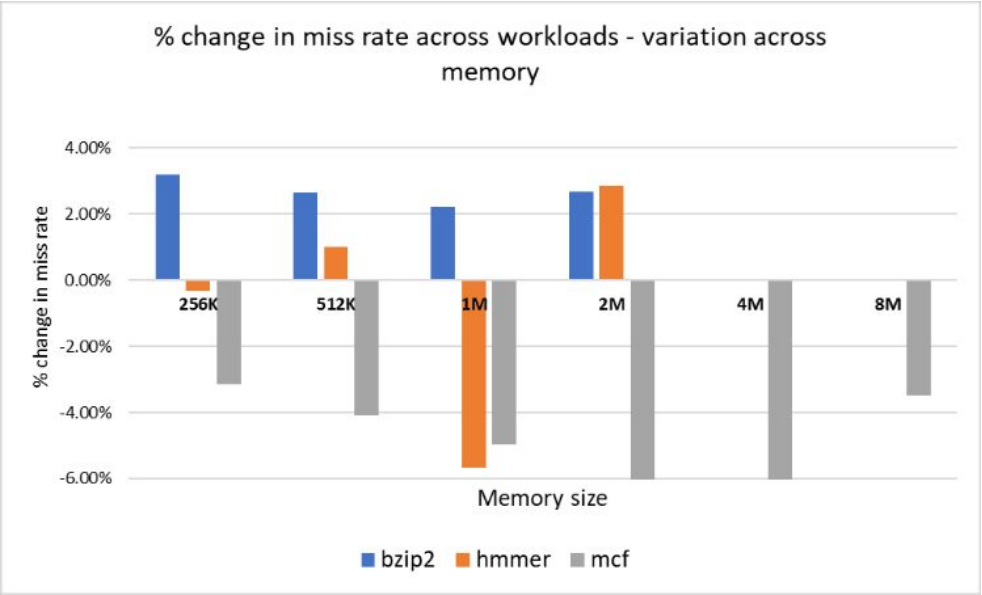

Effect of prefetcher

- As expected the prefetcher prevents compulsory misses.
- So for BZIP and HMMR, the prefetcher decreases the miss rate for both LRU and SRRIP.
- MCF doesn't have any change, which might mean that compulsory misses could not predicted by the prefetcher.

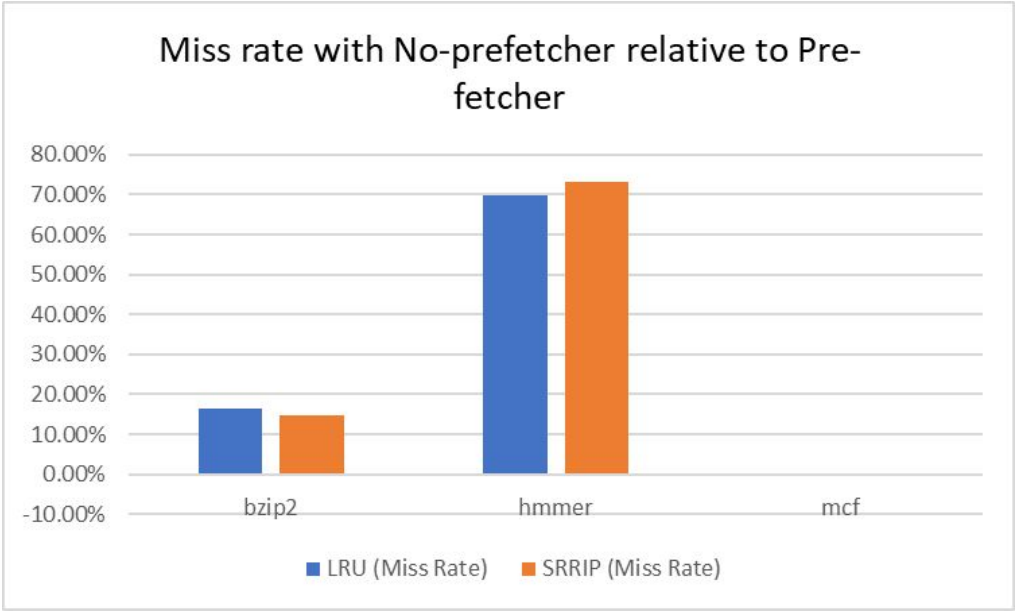

S-curve

- For the 30 datapoints we analysed, most of the datapoints that show negative performance benefit are from BZIP.
- HMMER shows zero to little benefit, while MCF shows the highest benefit with SRRIP policy as compared to RRIP.

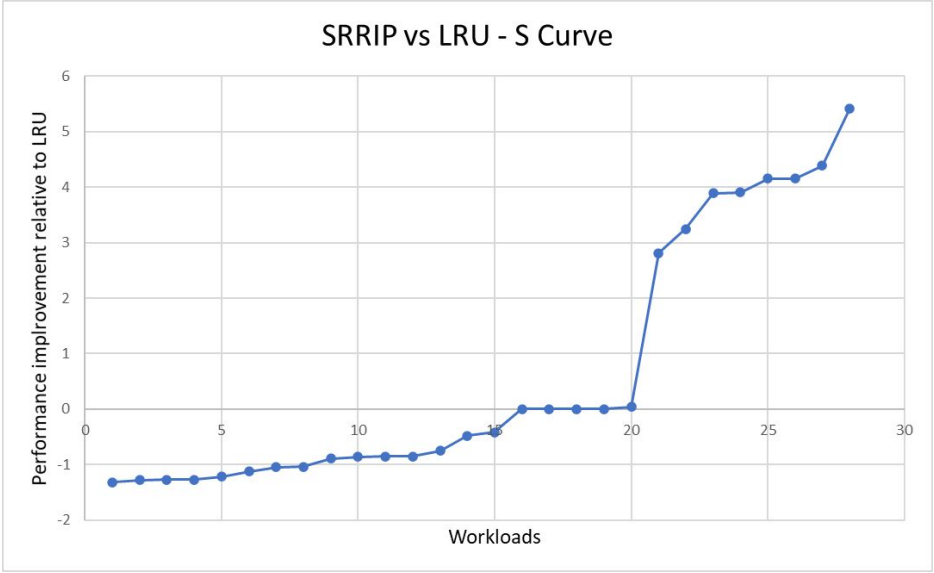

Learnings

- Cache replacement policy is affected by -
  - Workload- Relative size of workload to cache size
  - Memory access pattern in a workload
  - Associativity

References

- https://www.spec.org/cpu2006/.
- Nair AA, John LK, "Simulation points for SPECCPU 2006", Computer Design 2008, ICCD 2008, Pgs 397-403.
- A. Jaleel, K. Theobald, S. C. Steely, and J. Emer, “High Performance Cache Replacement Using Re-Reference Interval Prediction (RRIP),” in ISCA-32, 2010.